

CS221 Autumn 2021: Artificial Intelligence: Principles and Techniques

Homework 2: Sentiment Analysis

Name: Shaurya Goyal
Email: shaurya@kgpian.iitkgp.ac.in

By turning in this assignment, I agree by the Stanford honor code and declare that all of this is my own work.

Advice for this homework:

- Words are simply strings separated by whitespace. Note that words which only differ in capitalization are considered separate (e.g. great and Great are considered different words).
- You might find some useful functions in `util.py`. Have a look around in there before you start coding.

Before you get started, please read the **Assignments** section on the course website thoroughly.

Problem 1: Building intuition

Here are two reviews of *Perfect Blue*, from Rotten Tomatoes:

	Panos Kotzathanasis <i>Asian Movie Pulse</i>	 "Perfect Blue" is an artistic and technical masterpiece; however, what is of utmost importance is the fact that Satoshi Kon never deteriorated from the high standards he set here, in the first project that was entirely his own. Full Review	January 26, 2020
	Derek Smith <i>Cinematic Reflections</i>	 [An] nime thriller [that] often plays as an examination of identity and celebrity, but ultimately gets so lost in its own complex structure that it doesn't end up saying much at all. Full Review Original Score: 2/4	August 19, 2006

Rotten Tomatoes has classified these reviews as “positive” and “negative,” respectively, as indicated by the intact tomato on the top and the splatter on the bottom. In this assignment,

you will create a simple text classification system that can perform this task automatically. We'll warm up with the following set of four mini-reviews, each labeled positive (+1) or negative (−1):

1. (−1) pretty bad
2. (+1) good plot
3. (−1) not good
4. (+1) pretty scenery

Each review x is mapped onto a feature vector $\phi(x)$, which maps each word to the number of occurrences of that word in the review. For example, the first review maps to the (sparse) feature vector $\phi(x) = \{\text{pretty} : 1, \text{bad} : 1\}$. Recall the definition of the hinge loss:

$$\text{Loss}_{\text{hinge}}(x, y, \mathbf{w}) = \max\{0, 1 - \mathbf{w} \cdot \phi(x)y\},$$

where x is the review text, y is the correct label, \mathbf{w} is the weight vector.

- a. [2 points] Suppose we run stochastic gradient descent once for each of the 4 samples in the order given above, updating the weights according to

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} \text{Loss}_{\text{hinge}}(x, y, \mathbf{w}).$$

After the updates, what are the weights of the six words (“pretty”, “good”, “bad”, “plot”, “not”, “scenery”) that appear in the above reviews?

- Use $\eta = 0.1$ as the step size.
- Initialize $\mathbf{w} = [0, 0, 0, 0, 0, 0]$.
- The gradient $\nabla_{\mathbf{w}} \text{Loss}_{\text{hinge}}(x, y, \mathbf{w}) = 0$ when margin is exactly 1.

What we expect: A weight vector that contains a numerical value for each of the tokens in the reviews (“pretty”, “good”, “bad”, “plot”, “not”, “scenery”), **in this order**.

Your Solution: $[0, -0.2, -0.1, -0.1, -0.1, 0.1]$

b. [2 points] Given the following dataset of reviews:

1. (-1) bad
2. $(+1)$ good
3. $(+1)$ not bad
4. (-1) not good

Prove that no linear classifier using word features can get zero error on this dataset. Remember that this is a question about classifiers, not optimization algorithms; your proof should be true for any linear classifier, regardless of how the weights are learned. Propose a single additional feature for your dataset that we could augment the feature vector with that would fix this problem.

What we expect:

1. A short written proof (~ 3 -5 sentences).
2. A viable feature that would allow a linear classifier to have zero error on the dataset (classify all examples correctly).

Your Solution: Feature vector table (a column is one data point):

bad	1	0	1	0
good	0	1	0	1
not	0	0	1	1
Label	-1	1	1	-1

Setting the weights, "bad" should have a negative weight and "good" should have a positive weight as their individual labels are negative and positive respectively.

Hence the weight of "not" needs to be positive for "not bad" and negative for "not good" which is not possible

If these points are plotted in 3d space, it is observed that "bad" and "not good" lie on a plane and "good" and "not bad" lie on either side of the plane. Alternatively the plane can pass through "good" and "not bad" with "bad" and "not good" on either side. A single linear classifier cannot differentiate between such data without error.

An additional feature "not good" can solve this issue. It is 1 if the phrase matches the data point and 0 otherwise. This adds the row

not good	0	0	0	1
----------	---	---	---	---

 to the the feature vector table

Weights = $[-1, 1, 2, -4]$ corresponding to [bad, good, not, not good] with a decision boundary at 0 can solve this issue. ie classify as 1 if $\mathbf{w} \cdot \mathbf{x} > 0$ and -1 if $\mathbf{w} \cdot \mathbf{x} < 0$

Similarly "not bad" can also be taken

Problem 2: Predicting Movie Ratings

Suppose that we are now interested in predicting a numeric rating for movie reviews. We will use a non-linear predictor that takes a movie review x and returns $\sigma(\mathbf{w} \cdot \phi(x))$, where $\sigma(z) = (1 + e^{-z})^{-1}$ is the logistic function that squashes a real number to the range $(0, 1)$. For this problem, assume that the movie rating y is a real-valued variable in the range $[0, 1]$.

NOTE: Do not use math software such as Wolfram Alpha to solve this problem.

- a. [2 points] Suppose that we wish to use **squared loss**. Write out the expression for $\text{Loss}(x, y, \mathbf{w})$ for a single datapoint (x, y) .

What we expect: A mathematical expression for the loss. Feel free to use σ in the expression.

Your Solution: $\text{Loss}(x, y, \mathbf{w}) = (\sigma(\mathbf{w} \cdot \phi(x)) - y)^2$

- b. [3 points] Given $\text{Loss}(x, y, \mathbf{w})$ from the previous part, compute the gradient of the loss with respect to \mathbf{w} , $\nabla_{\mathbf{w}} \text{Loss}(x, y, \mathbf{w})$. Write the answer in terms of the predicted value $p = \sigma(\mathbf{w} \cdot \phi(x))$.

What we expect: A mathematical expression for the gradient of the loss.

Your Solution: $2(p - y) \frac{p^2(1-p)}{p} \phi(x) = \boxed{2(p - y)p(1 - p)\phi(x)}$

- c. [3 points] Suppose there is one datapoint (x, y) with some arbitrary $\phi(x)$ and $y = 1$. Specify conditions for \mathbf{w} to make the magnitude of the gradient of the loss with respect to \mathbf{w} arbitrarily small (i.e. minimize the magnitude of the gradient). Can the magnitude of the gradient with respect to \mathbf{w} ever be exactly zero? You are allowed to make the magnitude of \mathbf{w} arbitrarily large but not infinity.

What we expect:

1. 1-2 sentences describing the conditions for \mathbf{w} to minimize the magnitude of the gradient
2. 1-2 sentences explaining whether the gradient can be exactly zero.

HINT: Try to understand intuitively what is going on and what each part of the expression contributes. If you find yourself doing too much algebra, you're probably doing something suboptimal.

MOTIVATION: the reason why we're interested in the magnitude of the gradients is because it governs how far gradient descent will step. For example, if the gradient is close to zero when \mathbf{w} is very far from the optimum, then it could take a long time for gradient descent to reach the optimum (if at all). This is known as the *vanishing gradient problem* when training neural networks.

Your Solution: If magnitude of \mathbf{w} can be arbitrarily large but not infinity, the sigmoid function (or p) can take values very close to 0 or 1 (but not exactly) depending on the sign of $\mathbf{w} \cdot \phi(x)$. Since the gradient has both $(1 - p)$ and p terms, it doesn't matter if p tends to 0 or 1. As long as increasing the magnitude of \mathbf{w} increases the magnitude of $\mathbf{w} \cdot \phi(x)$, the gradient can be made arbitrarily small.

The gradient can be made 0 if $p = 0$ or 1. This is only possible in the limiting case where $\mathbf{w} \cdot \phi(x)$ is ∞ or $-\infty$. Since the magnitude of \mathbf{w} can't be infinity, it is only possible if $\phi(x)$ is of the form that can achieve infinity.

Problem 3: Sentiment Classification

In this problem, we will build a binary linear classifier that reads movie reviews and guesses whether they are “positive” or “negative.”

NOTE: Do not import any outside libraries (e.g. numpy) for any of the coding parts. Only standard python libraries and/or the libraries imported in the starter code are allowed. In this problem, you must implement the functions without using libraries like Scikit-learn.



- [2 points] Implement the function `extractWordFeatures`, which takes a review (string) as input and returns a feature vector $\phi(x)$ (you should represent the vector $\phi(x)$ as a `dict` in Python).
- [4 points] Implement the function `learnPredictor` using stochastic gradient descent and minimize the hinge loss. Print the training error and validation error after each epoch to make sure your code is working. You must get less than 4% error rate on the training set and less than 30% error rate on the validation set to get full credit.
- [2 points] Write the `generateExample` function (nested in the `generateDataset` function) to generate artificial data samples.

Use this to double check that your `learnPredictor` works! You can do this by using `generateDataset()` to generate training and validation examples. You can then pass in these examples as `trainExamples` and `validationExamples` respectively to `learnPredictor` with the identity function `lambda x:x` as a `featureExtractor`.

- [2 points] Some languages are written without spaces between words. So is splitting the words really necessary or can we just naively consider strings of characters that stretch across words? Implement the function `extractCharacterFeatures` (by filling in the `extract` function), which maps each string of n characters to the number of times it occurs, ignoring whitespace (spaces and tabs).
- [3 points] Run your linear predictor with feature extractor `extractCharacterFeatures`. Experiment with different values of n to see which one produces the smallest validation error. You should observe that this error is nearly as small as that produced by word features. Why is this the case? Construct a review (one sentence max) in which character n -grams probably outperform word features, and briefly explain why this is so.

NOTE: There is a function in `submission.py` that will allow you add a test to `grader.py` to test different values of n . Remember to write your final written solution here.

What we expect:

1. A short paragraph (4-6 sentences). In the paragraph state which value of n produces the smallest validation error, why this is likely the value that produces the smallest error.
2. A one-sentence review and explanation for when character n -grams probably outperform word features.

Your Solution: $n = 7$ produces the smallest validation error. The validation error decreases from $n = 1$ to $n = 7$ and then it starts increasing again.

The data has a lot of big words in the reviews. Hence the phrases formed by considering 7 letters at a time must be the most prevalent in the validation and training data hence a match can be found and the error is lower.

Character n -grams will probably outperform word features if a positive review has a lot of double negatives or vice-versa. The individual weights would have different weights than what the character n gram weight would be as individually the words mean the opposite as compared to when they are used together.

Problem 4: Toxicity Classification and Maximum Group Loss

Recall that models trained (in the standard way) to minimize the average loss can work well on average but poorly on certain groups, and that we can mitigate this issue by minimizing the maximum group loss instead. In this problem, we will compare the average loss and maximum group loss objectives on a toy setting inspired by a problem with real-world toxicity classification models.

Toxicity classifiers are designed to assist in moderating online forums by predicting whether an online comment is toxic or not, so that comments predicted to be toxic can be flagged for humans to review [1]. Unfortunately, such models have been observed to be biased: non-toxic comments mentioning demographic identities often get misclassified as toxic (e.g., “I am a [demographic identity]”) [2]. These biases arise because toxic comments often mention and attack demographic identities, and as a result, models learn to *spuriously correlate* toxicity with the mention of these identities.

In this problem, we will study a toy setting that illustrates the spurious correlation problem: The input x is a comment (a string) made on an online forum; the label $y \in \{-1, 1\}$ is the toxicity of the comment ($y = 1$ is toxic, $y = -1$ is non-toxic); $d \in \{0, 1\}$ indicates if the text contains a word that refers to a demographic identity; and $t \in \{0, 1\}$ indicates whether the comment includes certain “toxic” words. The comment x is mapped onto the feature vector $\phi(x) = [1, d, t]$ where 1 is the bias term (the bias term is present to prevent the edge case $\mathbf{w} \cdot \phi(x) = 0$ in the questions that follow). To make this concrete, we provide a few simple examples below, where we underline toxic words and words that refer to a demographic identity:

Comment (x)	Toxicity (y)	Presence of demographic mentions (d)	Presence of toxic words (t)
“Stanford <u>sucks</u> !”	1	0	1
“I’m a <u>woman</u> in computer science!”	-1	1	0
“The hummingbird <u>sucks</u> nectar from the flower”	-1	0	1

Suppose we are given the following training data, where we list the number of times each combination (y, d, t) shows up in the training set.

y	d	t	num. examples
-1	0	0	63
-1	0	1	27
-1	1	0	7
-1	1	1	3
1	0	0	3
1	0	1	7
1	1	0	27
1	1	1	63
Total num. examples			200

From the above table, we can see that 70 out of the 100 of toxic comments include toxic words, and 70 out of the 100 non-toxic comments do not. In addition, the toxicity of the comment t is highly correlated with mentions of demographic identities d (because toxic comments tend to target them) — 90 out of the 100 toxic comments include mentions of demographic identities, and 90 out of the 100 non-toxic comments do not.

We will consider linear classifiers of the form $f_{\mathbf{w}}(x) = \text{sign}(\mathbf{w} \cdot \phi(x))$, where $\phi(x)$ is defined above. Normally, we would train classifiers to minimize either the average loss or the maximum group loss, but for simplicity, we will compare two fixed classifiers (which might not minimize either objective):

- Classifier D: $\mathbf{w} = [-0.1, 1, 0]$
- Classifier T: $\mathbf{w} = [-0.1, 0, 1]$

For our loss function, we will be using the zero-one loss, so that the per-group loss is

$$\text{TrainLoss}_g(\mathbf{w}) = \frac{1}{|\mathcal{D}_{\text{train}}(g)|} \sum_{(x,y) \in \mathcal{D}_{\text{train}}(g)} \mathbf{1}[f_{\mathbf{w}}(x) \neq y].$$

Recall the definition of the maximum group loss:

$$\text{TrainLoss}_{\max}(\mathbf{w}) = \max_g \text{TrainLoss}_g(\mathbf{w}).$$

To capture the spurious correlation problem, let us define groups based on the value of (y, d) . There are thus four groups: $(y = 1, d = 1)$, $(y = 1, d = 0)$, $(y = -1, d = 1)$, and $(y = -1, d = 0)$. For example, the group $(y = -1, d = 1)$ refers to non-toxic comments with demographic mentions, for example.

- a. [2 points] In words, describe the behavior of Classifier D and Classifier T.

What we expect: For each classifier (D and T), an “if-and-only-if” statement describing the output of the classifier in terms of its features.

Your Solution:

Classifier D: Gives output +1 if and only if a demographic identity is mentioned, regardless of the presence of a toxic word. Gives -1 for all other

Classifier T: Gives output +1 if and only if a toxic word is present, regardless of the mention of a demographic identity. Gives -1 for all other

- b. [3 points] Compute the following three quantities concerning Classifier D using the dataset above:
1. Classifier D's average loss
 2. Classifier D's average loss for each group (fill in the table below)
 3. Classifier D's maximum group loss

What we expect: A value for average loss, a complete table (found below) with average loss for each group with the values in the given order, and a value for maximum group loss.

Classifier D	$y = 1$	$y = -1$
$d = 1$	0	1
$d = 0$	1	0

Your Solution:

1. Average Loss: $20/200 = 0.1$
2. In the table above
3. Maximum group loss: 1

c. [3 points] Now compute the following three quantities concerning Classifier T using the same dataset:

1. Classifier T's average loss
2. Classifier T's average loss for each group (fill in the table below)
3. Classifier T's maximum group loss

What we expect: A value for average loss, a complete table with average loss for each group with the values in the given order, and a value for maximum group loss.

Classifier T	$y = 1$	$y = -1$
$d = 1$	0.3	0.3
$d = 0$	0.3	0.3

Your Solution:

1. Average Loss: $60/200 = 0.3$
2. In the table above
3. Maximum group loss: 0.3

- d. [2 points] Now let's compare the two classifiers. Which classifier has lower average loss?
Which classifier has lower maximum group loss?

What we expect: First, indicate which classifier minimizes the average loss, then indicate which classifier minimizes the maximum group loss.

Your Solution:

Classifier D has lower average loss

Classifier T has lower maximum group loss

Since the maximum loss is minimized,

Classifier D minimizes the maximum group loss

Classifier T minimizes the average loss

- e. [2 points] As we saw above, different classifiers lead to different numbers of accurate predictions and different people's comments being wrongly rejected. Accurate classification of a non-toxic comment is good for the commenter, but when no classifier has perfect accuracy, how should the correct classifications be distributed across commenters? Here are four well-known principles of fair distribution:
- (a) According to **utilitarianism**, we should choose the distribution of accurate classifications that results in the *greatest net benefit* or *greatest average well-being*, where the average is a simple average.¹
 - (b) According to **prioritarianism**, we should choose the distribution of accurate classifications that results in the greatest *weighted* average well-being, where the weights prioritize less-well-off groups.²
 - (c) According to John Rawls's **difference principle**, when choosing between distributive systems, we should choose the one that maximizes the well-being of the worst-off people.³
 - (d) In order to **avoid compounding prior injustice**, we should ensure that our classifier does not impose a disadvantage on members of a demographic group that has faced historical discrimination.⁴

NOTE: The above are only a subset of the ethical frameworks out there. If you feel that another principle is more appropriate, you may invoke it, but you must define it in your answer and link to a source in which it is described. First, choose an objective (average loss or maximum group loss) by appealing to one of the four principles.

Which of Classifier D or Classifier T would you deploy on a real online social media platform in order to flag posts labeled as toxic for review? Refer to one of the above principles when justifying your answer.

What we expect: A 2-4 sentence explanation that justifies your choice of objective by referring to one of the four principles. There are many ways to answer this question well; a good answer explains the connection between a classifier and a principle clearly and concisely.

¹There are many varieties of utilitarianism and consequentialism – if you are interested in reading more about these principles, see here and here.

²For more on prioritarianism, see here and here.

³For more on the difference principle, see here and here.

⁴There are many anti-discrimination principles. This one is drawn from here; for others, see here. Many theories rely on perceived social group membership rather than demographic group membership. This principle also relates to corrective justice, for which see here.

Your Solution:

Assuming the training data distribution is representative of the actual distribution of comments:

There is a high number of toxic comments (45%) that are targeted at demographic identities ($d = 1$). Toxic comments that are not targetted to demographics are much less (5%). In this case, a Rawlsian approach would be more appropriate as demographic identities are worst off in comments due to the high number of comments with a demographic mention and a much higher probability of it being toxic. Classifier D would be chosen as it identifies correctly all toxic comments with a demographic identity term and non-toxic comments without a demographic term. It wrongly classifies toxic comments without a demographic term and non toxic comments with a demographic term but both combined just account for 10% of the total comments. These can be manually rectified. Classifying a non toxic comment as toxic is not as bad as classifying a toxic comment as non-toxic and since both are 5%, only 5% of the toxic comments are not being identified.

f. [2 points] We've talked about machine learning as the process of turning data into models, but where does the data come from? In the context of collecting data for training a machine learning model for toxicity classification, who determines whether a comment **should** be marked as toxic or not is very important (i.e., whether $y = 1$ or $y = -1$ in the earlier data table). Here are some commonly used choices:

- Recruit people on a crowdsourcing platform (like Amazon Mechanical Turk) to annotate each comment.
- Hire social scientists to establish criteria for toxicity and annotate each comment.
- Ask users of the platform to rate comments.
- Ask users of the platform to deliberate about and decide on community standards and criteria for toxicity, perhaps using a process of participatory design.⁵

Which method(s) would you use to determine the toxicity of comments for use as data to use for training a toxicity classifier, and why? Explain why you chose your method(s) over the others listed.

What we expect: 1-2 sentences explaining what methods(s) you would use and why you chose those method(s), and 1-2 sentences contrasting your chosen method(s) with alternatives.

Your Solution:

I would choose to hire social scientists as they have been trained to be objective and not let their personal biases affect their decisions and consider all possibilities of social implications. It is also highly likely that they have been exposed to confounding revelations about society and people in the course of their training so their criteria and annotation should be more inclusive and correct than the other choices.

In the other methods, there would be groups of people forming with similar biases to certain demographics etc and there is no guarantee that their ratings or deliberations would be free of their biases. There would be echo chambers and these biases would only amplify and many subtle toxic comments would not be annotated correctly. Asking the people of the platform where the toxic comments are coming from (which are 50% of all comments) is not a good choice as many of these people are themselves the source of these toxic comments

⁵For more on participatory design, see here

Problem 5: K-means Clustering

Suppose we have a feature extractor ϕ that produces 2-dimensional feature vectors, and a toy dataset $\mathcal{D}_{\text{train}} = \{x_1, x_2, x_3, x_4\}$ with

1. $\phi(x_1) = [10, 0]$
 2. $\phi(x_2) = [30, 0]$
 3. $\phi(x_3) = [10, 20]$
 4. $\phi(x_4) = [20, 20]$
- a. [2 points] Run 2-means on this dataset until convergence. Please show your work. What are the final cluster assignments z and cluster centers μ ? Run this algorithm twice with the following initial centers:
- (a) $\mu_1 = [20, 30]$ and $\mu_2 = [20, -10]$
 - (b) $\mu_1 = [0, 10]$ and $\mu_2 = [30, 20]$

What we expect: Show the cluster centers and assignments for each step.

Your Solution:

	Step	1	2	3
	Center 1	[20, 30]	[15, 20]	[15, 20]
	Center 2	[20, -10]	[20, 0]	[20, 0]
a)	x1 alloc	2	2	2
	x2 alloc	2	2	2
	x3 alloc	1	1	1
	x4 alloc	1	1	1

	Step	1	2	3
	Center 1	[0, 10]	[10, 10]	[10, 10]
	Center 2	[30, 20]	[25, 10]	[25, 10]
b)	x1 alloc	1	1	1
	x2 alloc	2	2	2
	x3 alloc	1	1	1
	x4 alloc	2	2	2

- b. [5 points] Implement the `kmeans` function. You should initialize your k cluster centers to random elements of `examples`. After a few iterations of k-means, your centers will be very dense vectors. In order for your code to run efficiently and to obtain full credit, you will need to precompute certain dot products. As a reference, our code runs in under a second on cardinal, on all test cases. You might find `generateClusteringExamples` in `util.py` useful for testing your code.

NOTE: Do not use libraries such as Scikit-learn.

- c. [2 points] If we scale all dimensions in our initial centroids and data points by some factor, are we guaranteed to retrieve the same clusters after running k-means (i.e. will the same data points belong to the same cluster before and after scaling)? What if we scale only certain dimensions? If your answer is yes, provide a short explanation; if not, give a counterexample.

What we expect: This response should have two parts. The first should be a yes/no response and explanation or counterexample for the first subquestion (scaling all dimensions). The second should be a yes/no response and explanation or counterexample for the second subquestion (scaling only certain dimensions).

Your Solution: No- Since K means defined in the question uses only euclidean distance which involves squaring a term, scaling all or certain dimensions does not provide a guarantee that the clusters after assignment will be same. Consider a data that is distributed from -100000000 to 100000000 on the x axis and -50 to 50 on the y axis. Due to the squared term, majority of the difference in distances from the center will be due to the x term. In this case, the algorithm may give clusters along the x axis (different clusters are horizontal). By scaling just the X axis, the cluster will change as now the error term will now be driven mainly by the y axis. By scaling the y axis also, there will be equal contribution of both the axes causing the clusters to change again.

Submission

Submission is done on Gradescope.

Written: When submitting the written parts, make sure to select **all** the pages that contain part of your answer for that problem, or else you will not get credit. To double check after submission, you can click on each problem link on the right side and it should show the pages that are selected for that problem.

Programming: After you submit, the autograder will take a few minutes to run. Check back after it runs to make sure that your submission succeeded. If your autograder crashes, you will receive a 0 on the programming part of the assignment. Note: the only file to be submitted to Gradescope is `submission.py`.

More details can be found in the Submission section on the course website.